

1 Import

```
In [1]: 1 %matplotlib inline
2 import warnings
3 warnings.filterwarnings("ignore")
4 import sqlite3
5 import pandas as pd
6 import numpy as np
7 import nltk
8 import string
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 from sklearn.feature_extraction.text import TfidfTransformer
12 from sklearn.feature_extraction.text import TfidfVectorizer
13 from sklearn.feature_extraction.text import CountVectorizer
14 from sklearn.metrics import confusion_matrix
15 from sklearn import metrics
16 from sklearn.metrics import roc_curve, auc
17 from nltk.stem.porter import PorterStemmer
18 import re
19 import string
20 from nltk.corpus import stopwords
21 from nltk.stem import PorterStemmer
22 from nltk.stem.wordnet import WordNetLemmatizer
23 from gensim.models import Word2Vec
24 from gensim.models import KeyedVectors
25 import pickle
26 import scipy as sp
27 from tqdm import tqdm
28 import os
29 from sklearn.cross_validation import train_test_split
30 from sklearn.neighbors import KNeighborsClassifier
31 from sklearn.metrics import accuracy_score
32 from sklearn.cross_validation import cross_val_score
33 from collections import Counter
34 from sklearn.metrics import accuracy_score
35 from sklearn import cross_validation
36 from sklearn.metrics import confusion_matrix
37 from sklearn.preprocessing import normalize
38 from sklearn import datasets, neighbors
39 from sklearn.metrics import roc_auc_score
40 from sklearn.preprocessing import StandardScaler
41
42 from sklearn.linear_model import LogisticRegression
43 from sklearn.model_selection import GridSearchCV
44 from sklearn.linear_model import SGDClassifier
45 from sklearn.svm import SVC
46 from sklearn import tree
47 from sklearn.cluster import KMeans, AgglomerativeClustering, DBSCAN
48 from sklearn.decomposition import TruncatedSVD
49
50 with open("C:/Python/Assignments/Preprocessing/final_sorted.txt", "rb") as file:
51     sorted_data = pickle.load(file)
```

C:\Anaconda3\lib\site-packages\gensim\utils.py:860: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")

C:\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
"This module will be removed in 0.20.", DeprecationWarning)

2 [5] Assignment 11: Truncated SVD

1. Apply Truncated-SVD on only this feature set:

- **SET 2:** Review text, preprocessed one converted into vectors using (TFIDF)
- **Procedure:**
 - Take top 2000 or 3000 features from tf-idf vectorizers using idf_score.
 - You need to calculate the co-occurrence matrix with the selected features (Note: X.X^T doesn't give the co-occurrence matrix, it returns the covariance matrix, check these blogs [blog-1](https://medium.com/data-science-group-jitir/word-embedding-2d05d270b285) (<https://medium.com/data-science-group-jitir/word-embedding-2d05d270b285>) [blog-2](https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2vec/) (<https://www.analyticsvidhya.com/blog/2017/06/word-embeddings-count-word2vec/>) for more information)
 - You should choose the n_components in truncated svd, with maximum explained variance. Please search on how to choose that and implement them. (hint: plot of cumulative explained variance ratio)
 - After you are done with the truncated svd, you can apply K-Means clustering and choose the best number of clusters based on elbow method.
 - Print out wordclouds for each cluster, similar to that in previous assignment.
 - You need to write a function that takes a word and returns the most similar words using cosine similarity between the vectors(vector: a row in the matrix after truncatedSVD)

2.1 Truncated-SVD

2.1.1 Load

```
In [2]: 1 # Taking 100k Points
2 my_final = sorted_data[:100000]
3 X = my_final['Text_new'].values
```

2.1.2 [5.1] Taking top features from TFIDF, SET 2

```
In [3]: 1 #Taking top 2000 features
2 tf_idf_vect = TfidfVectorizer( min_df=10, max_features=2000)
3 X_tfidf = tf_idf_vect.fit_transform(X)
```

```
In [4]: 1 # List of top 2000 words
2 top_words = np.array(tf_idf_vect.get_feature_names())
```

2.1.3 [5.2] Calculation of Co-occurrence matrix

```
In [5]: 1 sent = " ".join(X)
2 sent_list = sent.split(" ")
3 m, sent_length= len(top_words), len(sent_list)
4 dic = dict((top_words[i],i) for i in range(m))
5
```

```

In [7]: 1 import itertools
2 def create(neigh,m,sent_length,dic,sent_list):
3     '''
4     Parameters:
5     neigh = no of neighbours
6     m = length of top_words
7     sent_list = Corpus of all the words in data
8     Sent_length = length of corus
9     dic = Dictionary of each topwords and their location in topword list
10    '''
11    matrix = np.zeros((m,m)) # Initializing co occurrence matrix
12
13    tmp = sent_list[0:neigh+1] # The first window of neigh + 1 elements.
14    #print(tmp) # Since it is first we will check all possible combinations for matrix updation
15    for word in tmp: ## First update diag ele since it is same as the no of occurrence of the word
16        if(word in dic):
17            loc = dic[word] ## Get the Location in topword list
18            matrix[loc][loc] += 1
19    for w1,w2 in itertools.combinations(tmp, 2): # This will give all possible combinations (nC2)
20        if((w1 in dic) & (w2 in dic) & (w1 != w2)): # Bost must be top words and Both must be different
21            l1,l2 = dic[w1],dic[w2] # Then we get the Location and
22            #Update the matrix
23            matrix[l1][l2] += 1
24            matrix[l2][l1] += 1
25
26    #print("done 1st")
27    for i in tqdm(range(1,len(sent_list) - neigh)):
28        tmp = sent_list[i:i+neigh+1] # This is the window of neigh + 1 elements
29        #print(tmp)
30        new_word = tmp[neigh] # Only the Last added word is the new one and we have to pair this with other elements
31        if(new_word in dic): # only if it is a top word otherwise NO
32            loc = dic[new_word] # Location of the new word in topword list
33            matrix[loc][loc] += 1 # Because each diag ele is same as the no of occurrence of the word
34
35
36        for i in range(neigh): # pairing new word with other elements in the window
37            w1, w2 = tmp[i], new_word # getting the Location of both words in the topwords list
38            if((w1 in dic) & (w1 != new_word)): #if w1 is a topword and not same as new word, then we
39
40                l1,l2 = dic[w1], loc #get the Location in top word list
41                matrix[l1][l2] += 1 # AND dinaly update the matrix element
42                matrix[l2][l1] += 1
43
44    return matrix.astype(int)
45 co_matrix = create(neigh = 5, m = len(top_words), sent_length = len(sent_list),dic=dic, sent_list= sent_list)

```

100%|██| 3733264/3733264 [00:43<00:00, 86350.13it/s]

```

In [8]: 1 co_matrix

```

```

Out[8]: array([[3020, 1, 11, ..., 13, 9, 3],
 [ 1, 377, 0, ..., 3, 1, 1],
 [ 11, 0, 2616, ..., 3, 29, 13],
 ...,
 [ 13, 3, 3, ..., 958, 11, 3],
 [ 9, 1, 29, ..., 11, 2024, 3],
 [ 3, 1, 13, ..., 3, 3, 517]])

```

```

In [9]: 1 with open("mat.txt",'wb') as f:
2     pickle.dump(co_matrix,f)
3     # Saving this as a back up in case kernel is restarted

```

2.1.4 [5.3] Finding optimal value for number of components (n) to be retained.

```

In [10]: 1 matrix = co_matrix
2 svd = TruncatedSVD(n_components=2000-1)
3 data = svd.fit_transform(matrix)

```

```

In [11]: 1 cum_var = np.cumsum(svd.explained_variance_ratio_)

```

```


In [12]: 1 plt.plot(cum_var)
2 plt.grid(1)
3 plt.xlabel("n_components")
4 plt.ylabel("Cum Variance Ratio")
5 plt.title("Components vs Cum Var Plot")
6 plt.hlines(0.98,0,2000,colors='r') #A line to show 99% coverage
7 plt.show

```

```

Out[12]: <function matplotlib.pyplot.show(*args, **kw)>

```



```

In [13]: 1 np.where(cum_var>0.98)[0][0:5]

```

```

Out[13]: array([451, 452, 453, 454, 455], dtype=int64)

```

Observation: So with only 450 features I have 98% coverage of the whole data. So I will take 500 components instead of 2000

```

In [14]: 1 svd_new = TruncatedSVD(n_components=500)
2 data_new = svd_new.fit_transform(matrix)

```


2.1.5 [5.4] Applying k-means clustering

```

In [15]: 1 k_values = [2,3,4,5,6,7,8,9,10,11,12]
2 inertia = []
3 for k in (k_values):
4     kmeans = KMeans(n_clusters=k).fit(data_new)
5     inertia.append(kmeans.inertia_)
6

```

```
1 plt.plot(k_values,inertia,'r')
2 plt.grid(True)
3 plt.title("K vs Inertia Plot\n")
4 plt.xlabel("K Values")
5 plt.ylabel("Inertia Value")
6 plt.xticks(k_values)
7 plt.show();
```






Observation: At Kvalue = 4, there is an elbow in the graph, So I will take 4 clusters for KMeans

2.1.6 [5.5] Wordclouds of clusters obtained in the above section

```

1 opt_k = 4
2 kmeans = KMeans(n_clusters=opt_k).fit(data_new)
3 clust = [ [ i for i in range(opt_k) ] ] # this is the list of clusters
4 for i in range(kmeans.labels_.shape[0]):
5     clust[kmeans.labels_[i]].append(X[i])
6
7 from wordcloud import WordCloud
8 # https://www.geeksforgeeks.org/generating-word-cloud-python/
9 def cloud(cluster):
10     wordcloud = WordCloud(collocations=False, background_color = 'white',
11                           min_font_size = 10).generate(str(list(cluster)))
12     plt.figure(figsize = (6,6), facecolor = None)
13     plt.imshow(wordcloud)
14     plt.axis("off")
15     plt.tight_layout(pad = 0)
16     plt.title("Cluster",size = 30,)
17     plt.show()
18     for cl in clust:
19         cloud(cl)

```



Observation:

1. Cluster 1 is about tea, flavor, chocolate, Product related items.
2. Cluster 2 is about sauce, tabasco, mexican foods etc.

4. Cluster 4 is about beer, british food, tea, root etc.

▼

In [18]:

Out[18]: (2000, 2000)

In [19]:

In [22]:

```
Out[22]: ['across', 'came', 'come', 'upon', 'running']
```

2

3. I think I will get better and more clear clusters if I take more data points

window_size: 2

abc pqr def

abc	3.0	3.0	3.0
-----	-----	-----	-----

```
pqr      3.0      4.0      2.0
```

```
def      3.0      2.0      2.0
```

In [23]:

```
array([[3, 3, 3],
       [3, 4, 2],
       [3, 2, 2]])
```

Now, let us take an example corpus to calculate a co-occurrence matrix.

Corpus = He is not lazy. He is intelligent. He is smart.

	He	is	not	lazy	intelligent	smart
He	0	4	2	1	2	1
is	4	0	1	2	2	1
not	2	1	0	1	0	0
lazy	1	2	1	0	0	0
intelligent	2	2	0	0	0	0
smart	1	1	0	0	0	0

```
In [24]: 1 X1 = ['He is not lazy','He is intelligent','He is smart']
2 top_words1 = ['He','is','not','lazy','intelligent','smart']
3 sent1 = " ".join(X1)
4 sent_list1 = sent1.split(" ")
5 m = len(top_words1)
6 d1c1 = dict((top_words1[i],i) for i in range(m))
7 co_matrix = create(neigh = 2, m = len(top_words1), sent_length = len(sent_list1),dic=d1c1, sent_list= sent_list1)
8 co_matrix
```

```
100%|███████████████████████| 7/7 [00:00<?, ?it/s]
```

```
Out[24]: array([[3, 4, 2, 1, 2, 1],
                [4, 3, 1, 2, 2, 1],
                [2, 1, 1, 1, 0, 0],
                [1, 2, 1, 1, 0, 0],
                [2, 2, 0, 0, 1, 0],
                [1, 1, 0, 0, 0, 1]])
```